

Computer Graphics

5. GPU and Shaders

I-Chen Lin

National Yang Ming Chiao Tung University


Textbook: E. Angel, D. Shreiner Interactive Computer Graphics, 6th Ed., Pearson
Ref: D.D. Hearn, M. P. Baker, W. Carithers, Computer Graphics with OpenGL, 4th Ed., Pearson

Intended Learning Outcomes

- ▶ On completion of this chapter, a student will be able to:
 - ▶ Describe the concept of programmable graphics pipeline (compared to fixed-function graphics pipeline).
 - ▶ Identify the characteristics of three primary shaders: vertex, geometry, and fragment shaders.
 - ▶ Apply the essential transformations and shading algorithms with OpenGL shading language.

The Development of Graphics Cards (consumer-level): Early 90's

- ▶ VGA cards in the early 90's
 - ▶ Just output designated “bitmap”.
 - ▶ Some with 2D acceleration, ex. “Bitblt”
 - ▶ Ex. S3
- ▶ Interactive 3D(or 2.5D) games relied on software rendering.
 - ▶ There were hardware graphics pipelines on workstations, e.g. SGI.



The Development of Graphics Cards (consumer-level): Late 90's

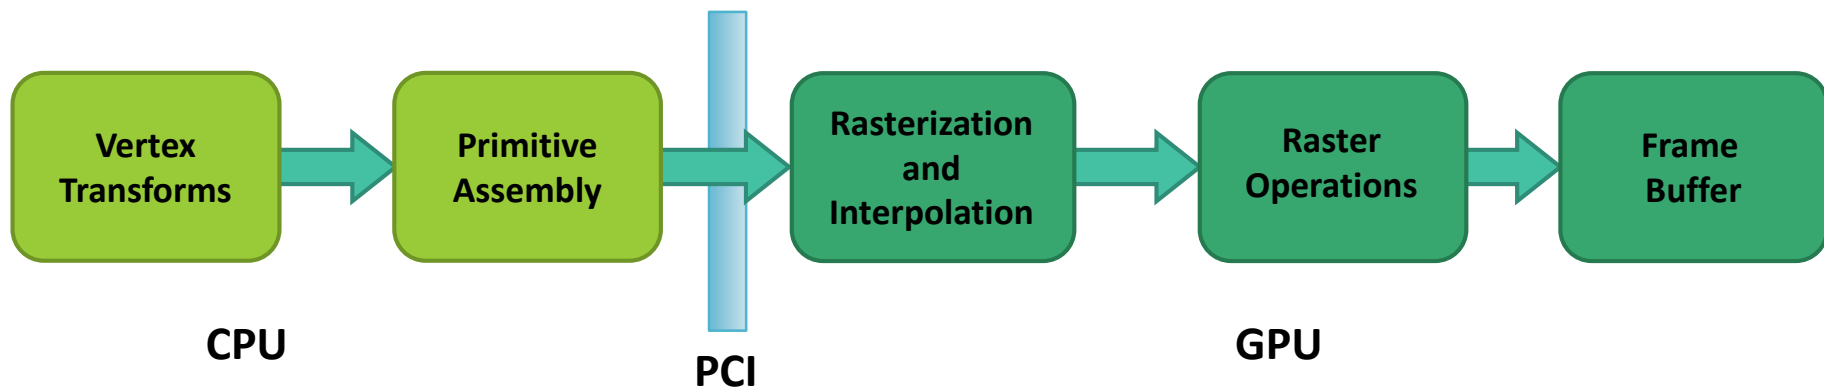
- ▶ 3D accelerators (90's)
 - ▶ Fixed-function pipelines.
 - ▶ E.g. S3, Voodoo, Nvidia, ATI, 3D Labs....
 - ▶ Some of them had to work with a standard VGA card.

3Dfx Voodoo (1996)

- ▶ One of the first true 3D game cards
- ▶ Worked by supplementing a standard 2D video card.
- ▶ Did not do vertex transformations (they were evaluated in the CPU)
- ▶ Did texture mapping, z-buffering.



en.wikipedia.org/wiki/3dfx_Interactive




Modified from S. Venkatasubramanian and J. Kider, “Evolution of the Programmable Graphics Pipeline”

GeForce/Radeon 7500 (1998)

- ▶ Main innovation: shifting the transformation and lighting calculations to the GPU
- ▶ Allowed multi-texturing: giving bump maps, light maps, and others.
- ▶ Faster AGP bus instead of PCI



en.wikipedia.org/wiki/GeForce_256



The Development of Graphics Cards (consumer-level): after 2001

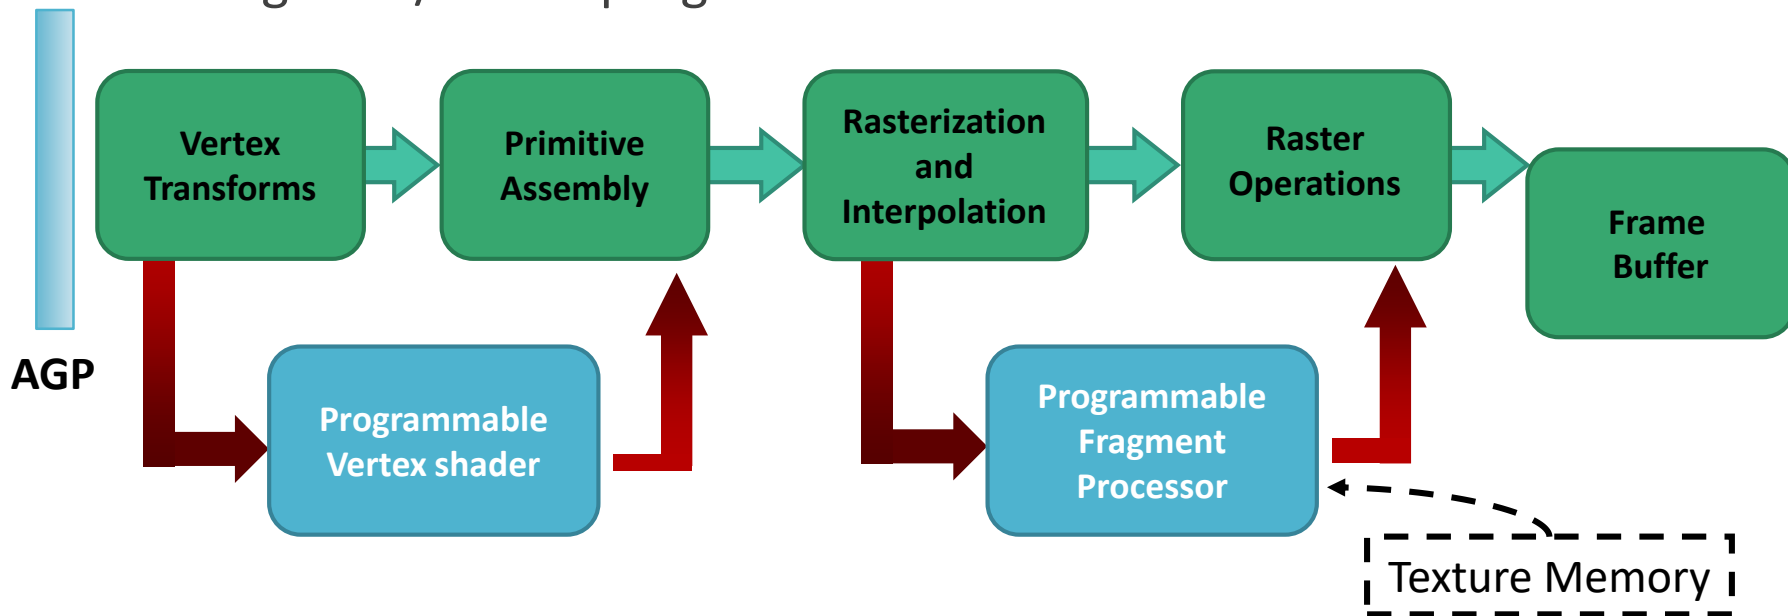
- ▶ Programmable pipelines on GPU
- ▶ GeForce3/Radeon 8500(2001)
 - ▶ Programmable vertex computations: up to 128 instructions
 - ▶ Limited programmable fragment computations: 8-16 instructions



https://en.wikipedia.org/wiki/GeForce_3_series

The Development of Graphics Cards (consumer-level): after 2001 (cont.)

- ▶ Radeon 9700/GeForce FX (2002)
 - ▶ the first generation of fully-programmable graphics cards
 - ▶ Different versions have different resource limits on fragment/vertex programs



Modified from S. Venkatasubramanian and J. Kider, "Evolution of the Programmable Graphics Pipeline"

Evaluation of Graphics Pipeline

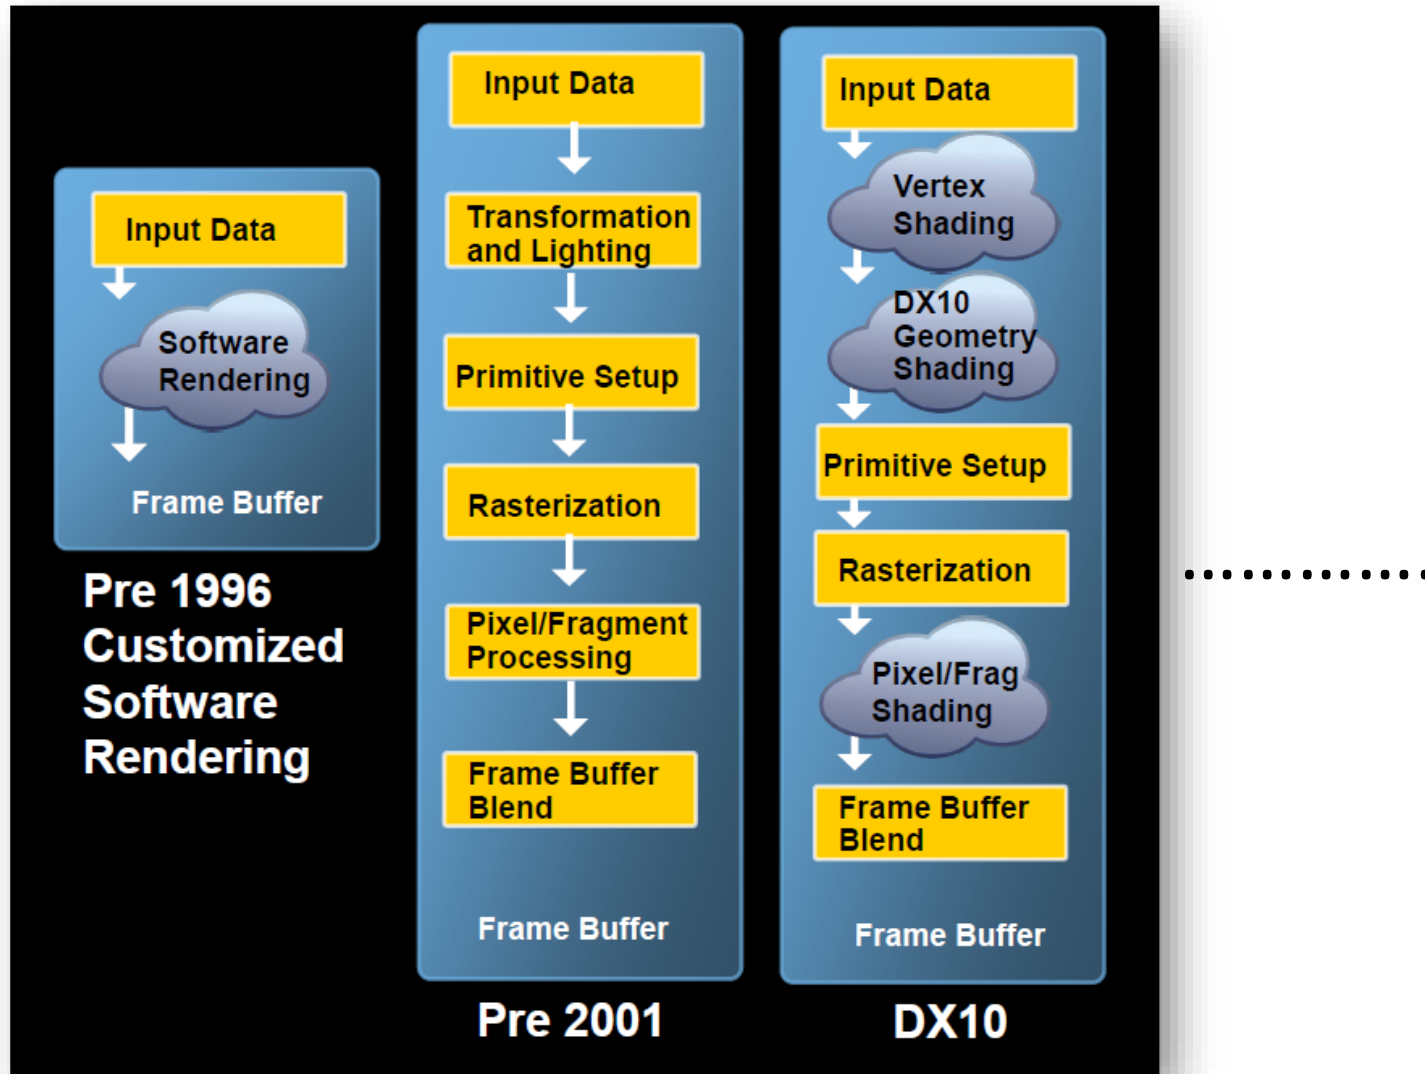


Figure from: M. Houston, "Beyond Programmable Shading Retrospective" slides

GPU & Shaders : the new age of real-time graphics

- ▶ Programmable pipelines.
- ▶ Supported by high-end commodity cards
 - ▶ NVIDIA, AMD/ATI, etc.



en.wikipedia.org/wiki/GeForce_10_series



www.amd.com/zh-hant/products/graphics/radeon-rx-570

Why is It So Remarkable?

- ▶ We can do lots of cool stuff in real-time, without overworking the CPU.
 - ▶ Phong Shading
 - ▶ Bump Mapping
 - ▶ Particle Systems
 - ▶ Animation
 - ▶
- ▶ Beyond real-time graphics: GP-GPU, e.g. CUDA, OpenCL (Open Computing Language)
 - ▶ Scientific data processing
 - ▶ Computer vision
 - ▶ Deep learning
 - ▶

Programmable Components

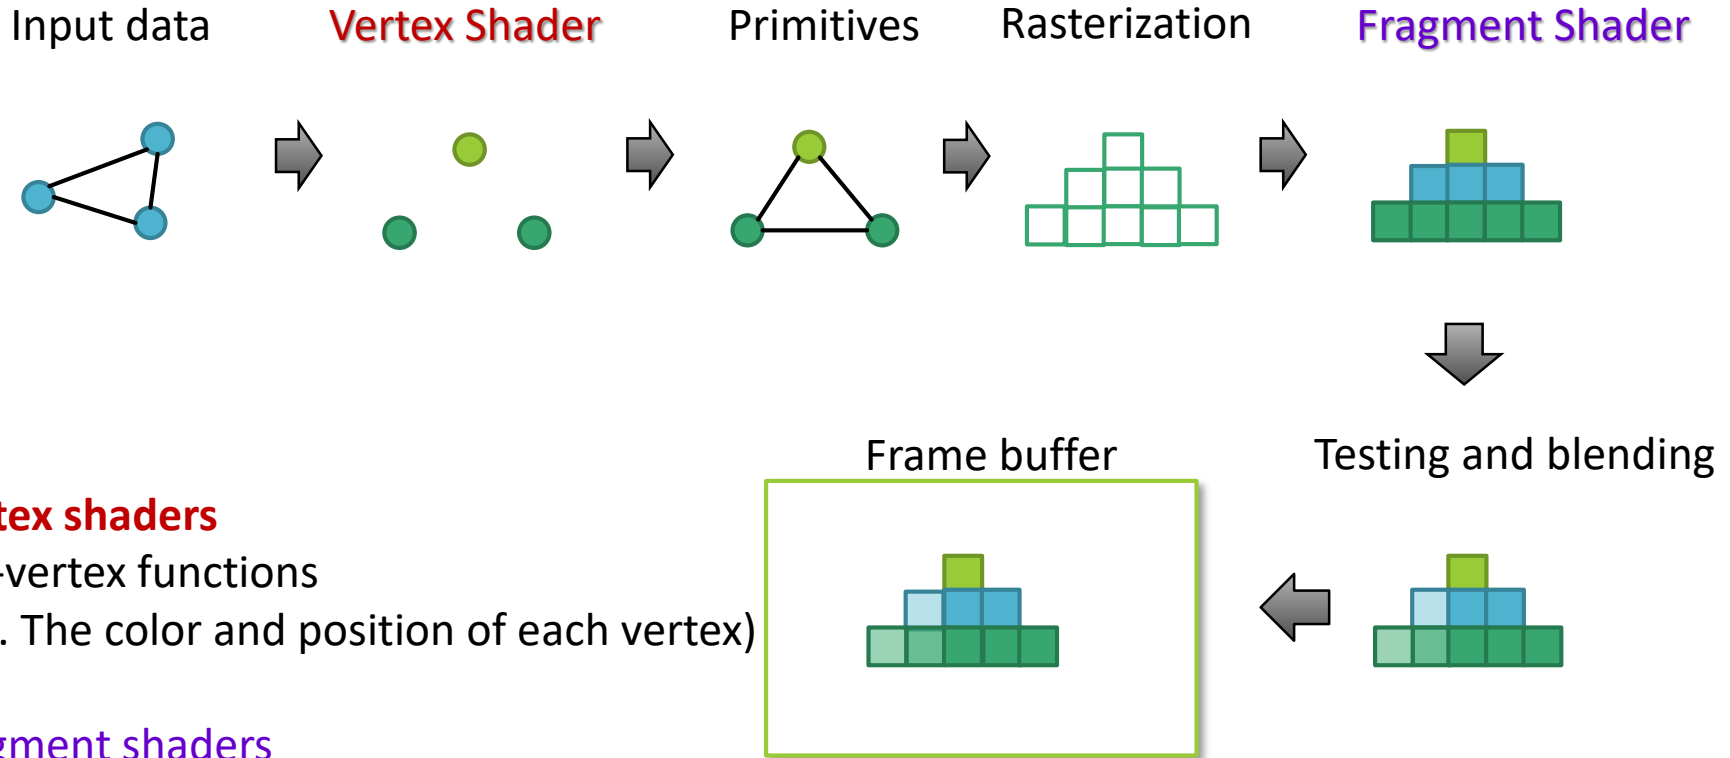
- ▶ **Shader:** programmable processors.
 - ▶ Replacing fixed-function vertex and fragment processing, and so forth.
- ▶ Types of shaders:
 - ▶ **Vertex shaders**
 - ▶ Dealing with per-vertex functions.
 - ▶ We can control the lighting and position of each vertex.
 - ▶ **Fragment shaders**
 - ▶ Dealing with per-pixel functions.
 - ▶ We can control the color of each pixel by user-defined programs.
 - ▶ **Geometry shaders** (DirectX 10, SM 4+)
 - ▶ New shaders (hull, domain) in DirectX11, SM5
 - ▶

Programmable Components (cont.)

▶ Software Support

- ▶ Direct X 8 , 9, 10, 11, 12, 12 Ultimate...
- ▶ OpenGL Extensions
- ▶ OpenGL Shading Language (GLSL)
- ▶ OpenGL for Embedded Systems (OpenGL ES)
- ▶ Vulkan
- ▶ Cg (*C for Graphics*)
- ▶ Metal Shading Language (by Apple)
- ▶

Essential GLSL pipeline (Vert.+Frag. Shaders)



Vertex shaders

per-vertex functions
(E.g. The color and position of each vertex)

Fragment shaders

per-fragment (pixel) function.
(E.g. The color of each fragment)

What about GLSL programs?

- ▶ Besides your main program (e.g. main.cpp), there are additional shader codes.
- ▶ These code can be character strings in your .cpp, but we usually put them in separate files (e.g. ooo.vs or ooo.vert, xxx.fs or xxx.frag).
- ▶ In a GLSL program, you can use multiple (different) shader codes to demonstrate different illumination algorithms for objects or regions.

*Vert buffer passed
from main.cpp*

A simple example of shader codes

Vertex shader code

```
#version 330 core
layout (location = 0) in vec3 aPos; // the position variable has attribute position 0
```

```
out vec4 vertexColor; // specify a color output to the fragment shader
```


```
void main()
{
    gl_Position = vec4(aPos, 1.0); // see how we directly give a vec3 to vec4's constructor
    vertexColor = vec4(0.5, 0.0, 0.0, 1.0); // set the output variable to a dark-red color
}
```

Fragment shader code

```
#version 330 core
out vec4 FragColor;
```

```
in vec4 vertexColor; // the input variable from the vertex shader
```

```
void main()
{
    FragColor = vertexColor;
}
```



Note: **gl_FragColor** is deprecated The example is modified from samples in learnopengl.com

Another simple example

Codes in main.cpp

```
// For the vertex shader
```

```
int vertexShader = glCreateShader(GL_VERTEX_SHADER);  
glShaderSource(vertexShader, 1, &vertexShaderSource, NULL);  
glCompileShader(vertexShader);  
  
.....
```

```
// For the fragment shader
```

```
int fragmentShader = glCreateShader(GL_FRAGMENT_SHADER);  
glShaderSource(fragmentShader, 1, &fragmentShaderSource, NULL);  
glCompileShader(fragmentShader);  
  
.....
```

```
// link the above shaders
```

```
int shaderProgram = glCreateProgram();  
glAttachShader(shaderProgram, vertexShader);  
glAttachShader(shaderProgram, fragmentShader);  
glLinkProgram(shaderProgram);  
  
.....
```

The example is modified from samples in learnopengl.com

Another simple example (cont.)


Codes in main.cpp

```
float vertices[] = {  
    // positions    // colors  
    0.5f, -0.5f, 0.0f, 1.0f, 0.0f, 0.0f, // bottom right  
    -0.5f, -0.5f, 0.0f, 0.0f, 1.0f, 0.0f, // bottom left  
    0.0f, 0.5f, 0.0f, 0.0f, 0.0f, 1.0f // top  
};
```

```
unsigned int VBO, VAO;  
glGenVertexArrays(1, &VAO);  
glGenBuffers(1, &VBO);  
glBindVertexArray(VAO);
```

```
glBindBuffer(GL_ARRAY_BUFFER, VBO);  
glBufferData(GL_ARRAY_BUFFER, sizeof(vertices), vertices, GL_STATIC_DRAW);
```

```
// position attribute  
glVertexAttribPointer(0, 3, GL_FLOAT, GL_FALSE, 6 * sizeof(float), (void*)0);  
glEnableVertexAttribArray(0);  
// color attribute  
glVertexAttribPointer(1, 3, GL_FLOAT, GL_FALSE, 6 * sizeof(float), (void*)(3 * sizeof(float)));  
glEnableVertexAttribArray(1);  
.....  
glUseProgram(shaderProgram);
```



Another simple example (cont.)

Vertex shader code

```
#version 330 core

layout (location = 0) in vec3 aPos; // the position variable at position 0
layout (location = 1) in vec3 aColor; // the color variable at position 1

out vec3 ourColor; // output a color to the fragment shader

void main()
{   gl_Position = vec4(aPos, 1.0);
    ourColor = aColor; // set ourColor to the input vertex color
}
```

Vert buffer passed from main.cpp

```
float vertices[] = {
    // positions      // colors
    0.5f, -0.5f, 0.0f, 1.0f, 0.0f, 0.0f,
    -0.5f, -0.5f, 0.0f, 0.0f, 1.0f, 0.0f,
    0.0f, 0.5f, 0.0f, 0.0f, 0.0f, 1.0f
};
```


Fragment shader code

```
#version 330 core

out vec4 FragColor;

in vec3 ourColor;

void main()
{   FragColor = vec4(ourColor, 1.0);
}
```



An example with matrix passing

```
float vertices[] = {  
    // positions      // colors  
    0.5f, -0.5f, -1.5f, 0.5f, 0.0f, 0.0f, // Mid right  
    -0.5f, -0.5f, -1.5f, 0.5f, 0.0f, 0.0f, // Mid left  
    0.0f, 0.5f, -1.5f, 0.5f, 0.0f, 0.0f, // top  
    -0.5f, -0.5f, -1.5f, 0.0f, 0.5f, 0.0f, // Mid left  
    0.5f, -0.5f, -1.5f, 0.0f, 0.5f, 0.0f, // Mid right  
    0.0f, -1.5f, -1.5f, 0.0f, 0.5f, 0.0f // bottom  
};
```

Codes in main.cpp

*Data passing through
uniform variables*

```
.....  
glClearColor(0.0f, 0.0f, 0.0f, 1.0f);  
glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);  
  
glm::mat4 projection_matrix = glm::perspective(glm::radians(60.0f), (float)SCR_WIDTH / (float)SCR_HEIGHT,  
0.1f, 10.0f);  
glUniformMatrix4fv(glGetUniformLocation(shaderProgram, "projection"), 1, GL_FALSE,  
glm::value_ptr(projection_matrix) );  
  
glm::mat4 model_matrix = glm::mat4(1.0f); // make sure to initialize matrix to identity matrix first  
model_matrix = glm::rotate(model_matrix, glm::radians(15.0f), glm::vec3(0.0f, 0.0f, 1.0f));  
glUniformMatrix4fv(glGetUniformLocation(shaderProgram, "model"), 1, GL_FALSE, &model_matrix[0][0] );  
  
glDrawArrays(GL_TRIANGLES, 0, 6);
```

An example with matrix passing (cont.)


```
#version 330 core
layout (location = 0) in vec3 aPos;
layout (location = 1) in vec3 aColor;
out vec3 ourColor;
uniform mat4 model;
uniform mat4 projection;
```

Vertex shader code

```
void main()
{
    gl_Position = projection * model * vec4(aPos, 1.0);
    ourColor = aColor;
}
```

Fragment shader code

```
#version 330 core
out vec4 FragColor;
in vec3 ourColor;
void main()
{
    FragColor = vec4(ourColor, 1.0f);
}
```



Vertex Shaders

- ▶ Per-vertex calculations performed here
 - ▶ Without knowledge about other vertices (parallelism)
- ▶ Your program take responsibility for:
 - ▶ Vertex transformation
 - ▶ Normal transformation
 - ▶ (Per-Vertex) Lighting
 - ▶ Color material application and color clamping
 - ▶ Texture coordinate generation

Vertex Shader Applications

- ▶ We can control movement with uniform variables and vertex attributes
 - ▶ Time
 - ▶ Velocity
 - ▶ Gravity
- ▶ Moving vertices
 - ▶ Morphing
 - ▶ Wave motion
 - ▶
- ▶ Lighting
 - ▶ More realistic models
 - ▶ Cartoon shaders

Applications: Wave Motion Vertex Shader

Uniform: passing parameters to vertex and fragment shaders.

```
.....  
uniform float time;
```

```
uniform float xs, zs;
```

<- uniform parameters to all vertices.

For vertex-independent parameters, we can use

```
void main()
```

```
{vec3 object_pos;
```

```
float s;
```

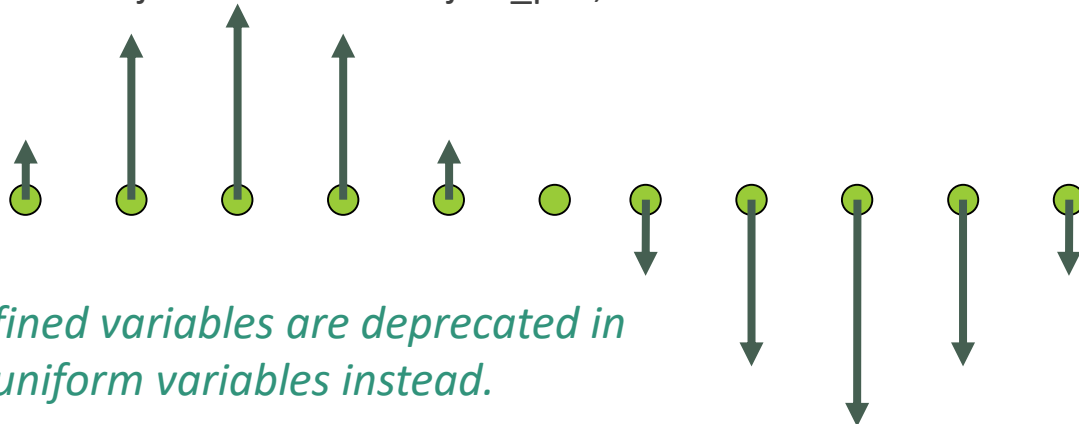
```
s = 1.0 + 0.1*sin(xs*time)*sin(zs*time);
```

```
object_pos = aPos;
```

```
object_pos.y = s* aPos.y;
```

```
gl_Position = gl_ModelViewProjectionMatrix* object_pos;
```

```
}
```





Note: Several `gl_` predefined variables are deprecated in the newer version. Use uniform variables instead.

Applications: Particle Systems

Uniform: passing parameters to vertex and fragment shaders.

.....

```
uniform vec3 vel;      <- uniform param. to all vertices.  
                        For vertex independent param.,.....  
uniform float g, t;  
  
void main()  
{  
    vec3 object_pos;  
    object_pos.x = aPos.x + vel.x*t;  
    object_pos.y = aPos.y + vel.y*t + g/(2.0)*t*t;  
    object_pos.z = aPos.z + vel.z*t;  
    gl_Position = gl_ModelViewProjectionMatrix*  
    vec4(object_pos,1);  
}
```




Note: Several gl_ predefined variables are deprecated in the newer version. Use uniform variables instead.

Fragment Shaders


- ▶ What is a fragment?
 - ▶ Cg Tutorial says: “You can think of a fragment as a ‘potential pixel’”
- ▶ Perform per-pixel calculations
 - ▶ Without knowledge about other fragments (parallelism)
- ▶ Your program’s responsibilities:
 - ▶ Operations on interpolated values
 - ▶ Texture access and application
 - ▶ Other functions: fog, color lookup, etc.

Fragment Shader Applications

(Per-pixel) Phong shading




Per-vertex lighting




Per-fragment lighting

Figures from <http://www.lighthouse3d.com/opengl/glsl/>

Fragment Shader Applications





smooth shading



bump mapping

Bump Mapping

- ▶ Perturb normal for each fragment
- ▶ Store perturbation as textures



Toon Shading

ftransform(): result from the GL fixed-function transformation pipeline

*Note: **varying**, communicating between vertex and fragment. Use **in out** variables in newer versions.*

- The vertex shader then becomes:

```
.....  
out vec3 vnormal;  
void main() {  
    vnormal = gl_NormalMatrix * gl_Normal;  
    gl_Position = ftransform(); }
```


- The fragment shader becomes

```
.....  
in vec3 vnormal;  
out vec4 FragColor;  
  
void main() {  
    float intensity; vec4 color;  
    vec3 n = normalize(vnormal);  
    intensity = dot(vec3(gl_LightSource[0].position),n);  
    if (intensity > 0.95) color = vec4(1.0,0.5,0.5,1.0);  
    else if (intensity > 0.5) color = vec4(0.6,0.3,0.3,1.0);  
    else if (intensity > 0.25) color = vec4(0.4,0.2,0.2,1.0);  
    else color = vec4(0.2,0.1,0.1,1.0);  
    FragColor = color; }
```



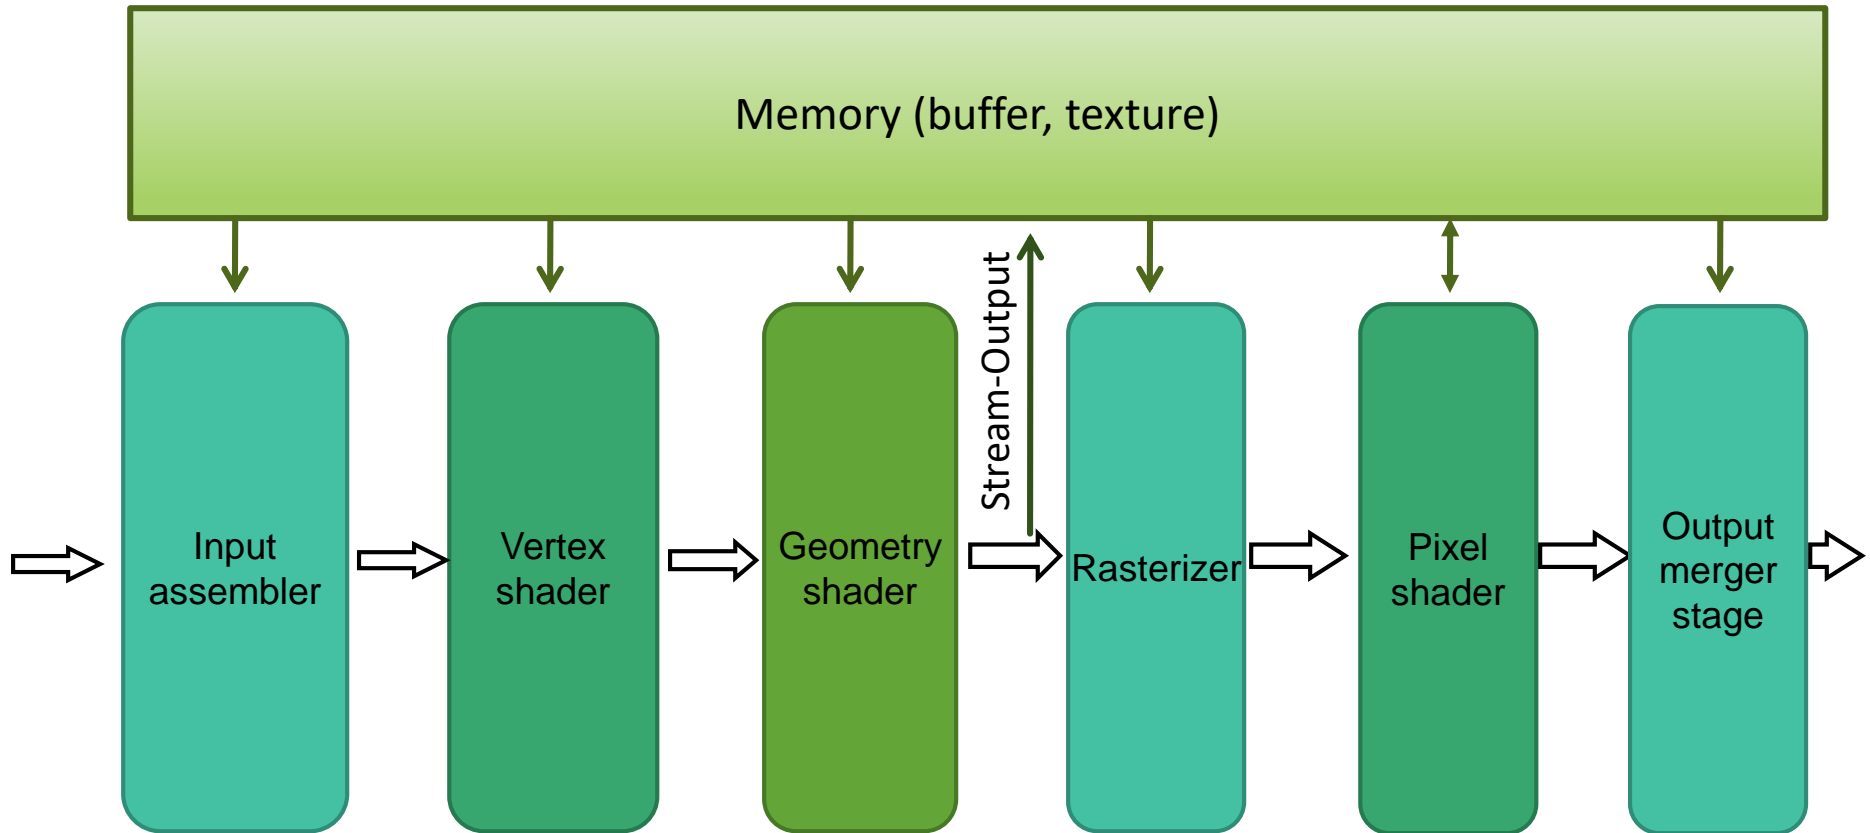
There are various ways to implement Toon shading

gl_NormalMatrix



- ▶ Can we directly apply the modelview matrix M to a normal vector ?
 - ▶ Problem: If the upper-left 3x3 submatrix M_s is not orthogonal, $n' = M_s n$ is not perpendicular to $v' = M_s v$

With the Geometry Shader



Direct3D 10 pipeline stage from MSDN of Microsoft


D3D 10 Pipeline

- ▶ **Input assembler:** supplies data (triangles, lines and points) to the pipeline.
- ▶ **Vertex shader:** processes vertices, such as transformations, skinning, and lighting.
- ▶ **Geometry shader:** processes entire primitives.
 - ▶ 3 vertices: a triangle, 2 vertices: a line, or 1 vertex: a point.
 - ▶ The Geometry shader supports limited geometry amplification and de-amplification. (discard the primitive, or emit one or more new primitives)
 - ▶ E.g. Subdivision, point -> billboard, silhouette edge -> fur, etc.
- ▶ **Stream-output stage:**
 - ▶ Data can be streamed out and/or passed into the rasterizer. Data streamed out to memory can be recirculated back into the pipeline as input data or read-back from the CPU.

D3D 10 Pipeline (cont.)

- ▶ **Rasterizer:** clips primitives, prepares primitives for the pixel shader and determines how to invoke pixel shaders.
- ▶ **Pixel shader:** receives interpolated data for a primitive and generates per-pixel data, such as color.
- ▶ **Output-merger stage:**
 - ▶ combines various types of output data (pixel shader values, depth and stencil information) with the contents of the render target and depth/stencil buffers to generate the final pipeline result.

GLSL pipeline (Vert.+Geo.+Frag. Shaders)




Fragment Shader



Testing and blending



Frame buffer

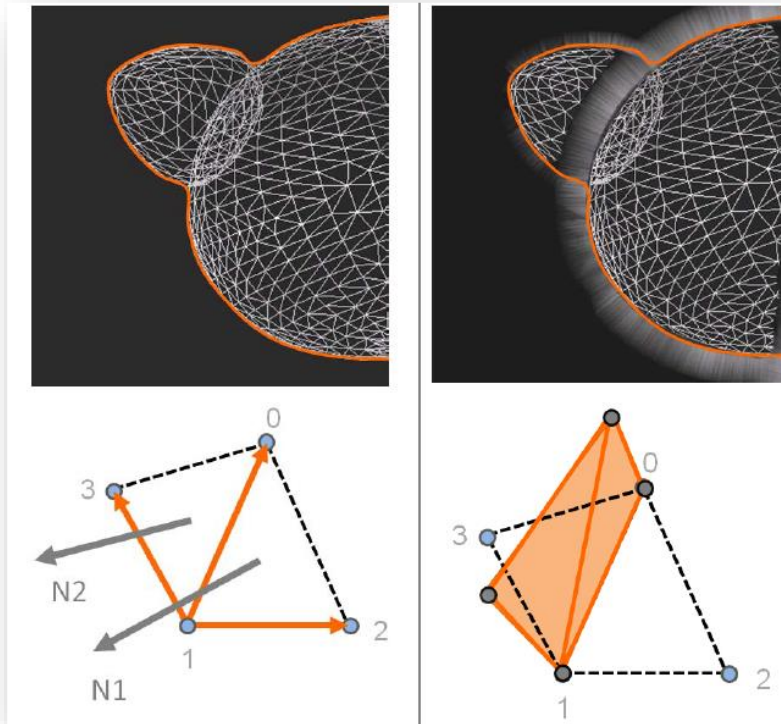


Vertex shaders
per-vertex functions


Geometry shaders
Primitive processing
(E.g. transformation, generating zero to multiple primitives)

Fragment shaders
per-fragment (pixel) function.

D3D 10 Pipeline (cont.)



Figures from NVIDIA DirectX10 SDK Doc:
Fur (using Shells and Fins)



Previous example using Geometry Shader

Vertex shader code

```
#version 330 core
layout (location = 0) in vec3 aPos;
layout (location = 1) in vec3 aColor;
out VS_OUT{
    vec3 vsColor;
} vs_out;
uniform mat4 model;
uniform mat4 projection;

void main()
{ gl_Position = projection * model * vec4(aPos, 1.0);
  vs_out.vsColor = aColor;
}
```

Fragment shader code


```
#version 330 core
out vec4 FragColor;
in vec3 ourColor;
void main()
{ FragColor = vec4(ourColor, 1.0f);
}
```

Geometry shader code

```
#version 330 core
layout (triangles) in;
layout (triangle_strip, max_vertices = 3) out;
in VS_OUT{
    vec3 vsColor;
} gs_in[];
out vec3 ourColor;
```

*Interface block names
e.g. VS_OUT
should be matched.*

```
void main()
{
    for(int i=0; i<3; i++)
    { gl_Position = gs_in[i].gl_Position;
      ourColor = gs_in[i].vsColor;
      EmitVertex();
    }
    EndPrimitive();
}
```



Adding additional triangles with GS

Geometry shader code

```
#version 330 core
layout (triangles) in;
layout (triangle_strip, max_vertices = 4) out;
in VS_OUT{
    vec3 vsColor;
} gs_in[];
out vec3 ourColor;
```

```
void main()
{ for(int i=0; i<3; i++)
  { gl_Position = gl_in[i].gl_Position;
    ourColor = gs_in[i].vsColor;
    EmitVertex();
  }
}
```

```
gl_Position = gl_in[0].gl_Position + vec4(0.2f, 0.2f, -0.2f, 0.0f);
ourColor = vec3(0.0f, 0.0f, 0.8f);
EmitVertex();
EndPrimitive();
```

```
}
```

Note:

Triangle strip: v0, v1, v2, v3

⇒ Triangle 1 (v0, v1, v2)

⇒ Triangle 2 (v1, v2, v3)

For each triangle, add one additional triangle.

The above code works, but

D3D 11 Pipeline

- ▶ In D3D10, the Geometry shader may subdivide the surfaces by multiple passes.
- ▶ D3D11 improves the tessellation ability by three new stages: hull shader, tessellator, domain shader.
- ▶ The tessellated patches can still be applied to geometry shaders. E.g. point -> billboard, silhouette edge -> fur, etc.

Tessellation Pipeline

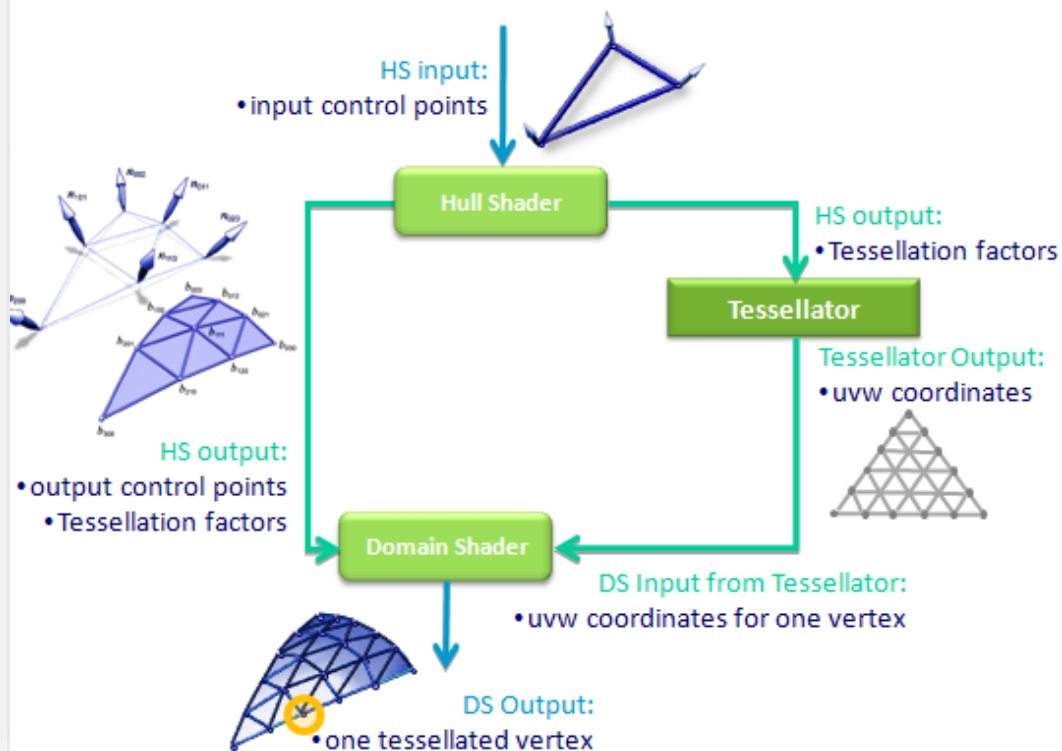


Figure from:
developer.download.nvidia.com/presentations/2009/GDC/GDC09_D3D11Tessellation.pdf





Figure from: vulkan-tutorial.com/Drawing_a_triangle/Graphics_pipeline_basics/Introduction

D3D 11 Tessellation



Model refinement



Base
Model




Bump
Mapping




Displacement
Mapping

Image courtesy of www.chromosphere.com



Tessellation with displacement mapping


Shading with Shaders: Examples



Toon shading (or cel-shading) with outline

- ▶ I1: The back faces are drawn with thick lines.
- ▶ I2: The object faces are drawn using a single color.
- ▶ I3: Toon shading is applied.

(Note: there are various methods to get outlines.)



Figures from en.wikipedia.org/wiki/Cel_shading

End of Chapter