

# **Computer Graphics**

## **8. Shadow**

I-Chen Lin

Institute of Multimedia Engineering,  
National Yang Ming Chiao Tung University

# Intended Learning Outcomes


- ▶ On completion of this chapter, a student will be able to:
  - ▶ Explain how and when the shadows occur.
  - ▶ Describe the typical algorithms for real-time shadow rendering.
  - ▶ Compare the advantages and limitations of typical shadow algorithms.
  - ▶ Apply the shadow mapping method with GLSL.

# Objective


- ▶ This chapter focuses on shadows in real-time graphics.
- ▶ Global illumination methods can generate realistic shadows but require heavy computation.
- ▶ Local illumination does not consider the situation where the light can be blocked by other objects.

# Why are shadows important?



- ▶ Depth cue
- ▶ Contact point




- ▶ Realism of illuminated scenes




# Why do shadows occur?



# Umbra and penumbra



Point light source



Area light source

# Real-time shadow techniques

- ▶ Projected planar shadows
  - ▶ works only on flat surfaces
- ▶ Light maps or texture shadows
  - ▶ unsuited for dynamic shadows
- ▶ Popular methods
  - ▶ Shadow map
  - ▶ Shadow volume


# Shadow mapping

- ▶ Image-space shadow
  - ▶ L. Williams published the idea in 1978.
- ▶ Completely image-space algorithm
  - ▶ No knowledge of scene's geometry is required
  - ▶ Aliasing artifacts may occur
- ▶ Nearly “standard” for shadow rendering
  - ▶ Pixar's RenderMan also include the algorithm
  - ▶ Applied in several Pixar's movies, e.g. Luxo Jr., Toy Story, etc.




# The concept of shadow mapping

- ▶ Depth testing from the light's point-of-view
  - ▶ Two-pass algorithm
    - ▶ First, render depth buffer from the light's point-of-view.
    - ▶ The result is a “shadow map”.
    - ▶ A 2D function indicating the depth of the closest pixels to the light
- ▶ This depth map is used in the second pass.




# The concept of shadow mapping (cont.)

- ▶ Second, render scene from the eye's point-of-view
- ▶ For each rasterized fragment
  - ▶ determine fragment's XYZ position *relative to the light*.
  - ▶ compare the depth value at light position XY in the depth map to fragment's light position Z.
- ▶ Comparing two values:
  - ▶  $A = Z$  value from depth map at fragment's XY position from the view of light.
  - ▶  $B = Z$  value of fragment with respect to light position.
  - ▶ If  $B$  is greater than  $A$ , then there must be something closer to the light than the fragment then the fragment is shadowed.
  - ▶ If  $A$  and  $B$  are approximately equal, the fragment is lit.




# Shadow mapping with a picture in 2D (1)

The  $A < B$  shadowed fragment case




# Shadow mapping with a picture in 2D (2)

The  $A \cong B$  unshadowed fragment case



# Shadow mapping with a picture in 2D (3)

Note image precision mismatch!




The depth map could be at a different resolution from the framebuffer

This mismatch can lead to artifacts


# Visualizing the shadow mapping (1)

- A fairly complex scene with shadows




# Visualizing the shadow mapping (2)

- Compare with and without shadows





with shadows



without shadows

# Visualizing the shadow mapping (3)

- The scene from the light's point-of-view





FYI: from the  
eye's point-of-view  
again



# Visualizing the shadow mapping (4)



- The depth buffer from the light's point-of-view



FYI: from the  
light's point-of-view  
again

# Visualizing the shadow mapping (5)


- Projecting the depth map onto the eye's view



FYI: depth map for  
light's point-of-view  
again


# Visualizing the shadow mapping (6)

- Projecting light-to-fragment distance onto eye's view



# Visualizing the shadow mapping (6)

- Comparing light distance to light depth map




Non-green is  
where shadows  
should be

# Visualizing the shadow mapping (7)

## ► Scene with shadows


*Notice how  
specular  
highlights never  
appear in  
shadows*



*Notice how  
curved surfaces  
cast shadows  
on each other*



# Shadow mapping problem (1)

- ▶ Accuracy of depth values (quantization and sampling)
  - ▶  $D_{sm}$  and  $D_{eye}$  may have different values even if they represent the same surface.
- ▶ Solution:  $D_{sm} + \text{bias} < D_{eye} \rightarrow \text{shadow}$



Too little bias, everything  
begins to shadow

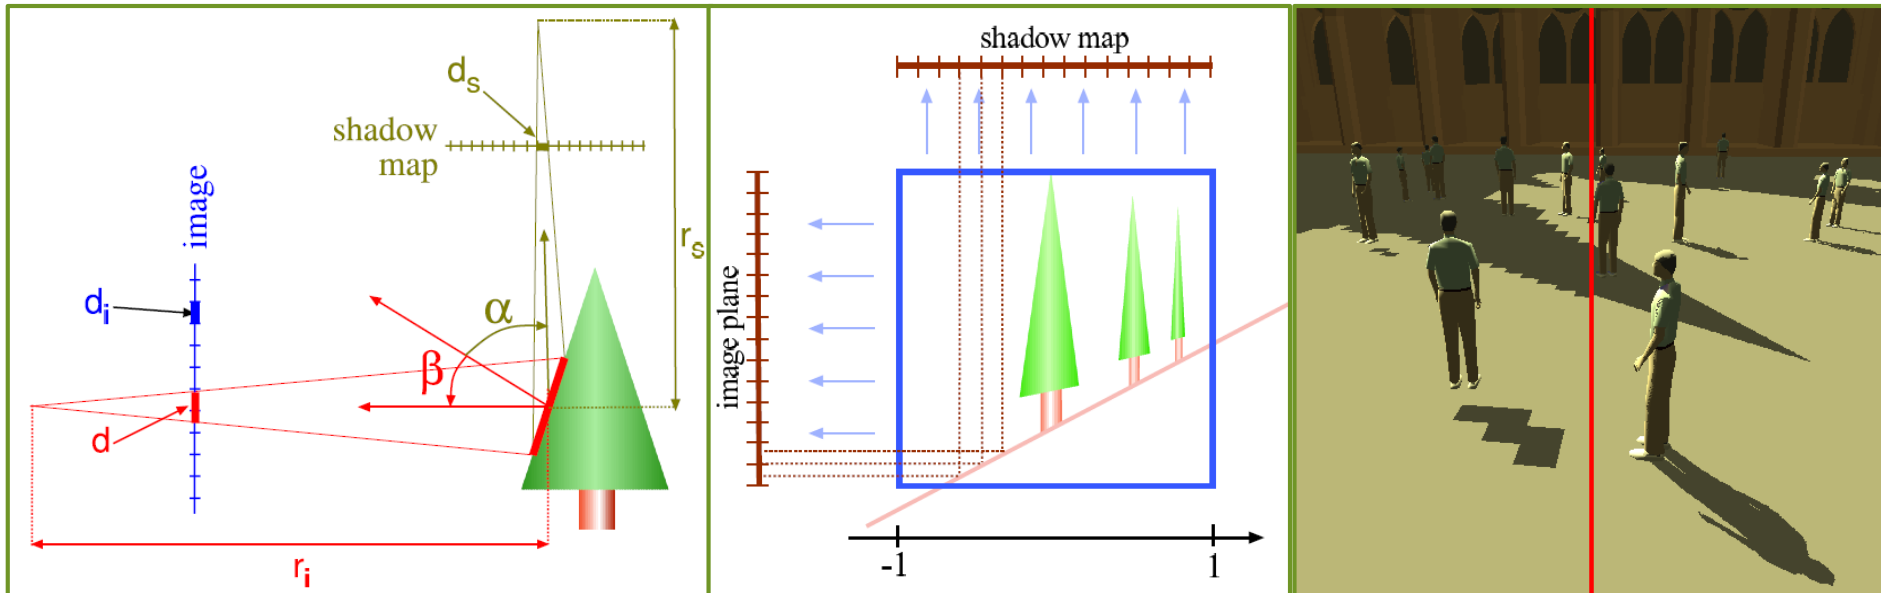
Just right



Too much bias, shadow  
starts too far back

# Shadow mapping problem (2)



- ▶ Insufficient resolution of a shadow map
  - ▶ Gives jagged shadow edges
  - ▶ Resolution needed depends on camera position
- ▶ Lots of research on improving quality



M. Stamminger, "Perspective Shadow Map", SIGGRAPH'02.

# “Percentage Closer” filtering

- ▶ Normal texture filtering just averages color components.
- ▶ Averaging depth values does NOT work.
- ▶ Provides anti-aliasing at shadow map edges.
  - ▶ Not the real soft shadows in the umbra/penumbra sense.
- ▶ Reeves et al. (pixar), Rendering antialiased shadows with depth maps, SIGGRAPH’87.






# Depth values are not blendable

- Traditional filtering is inappropriate

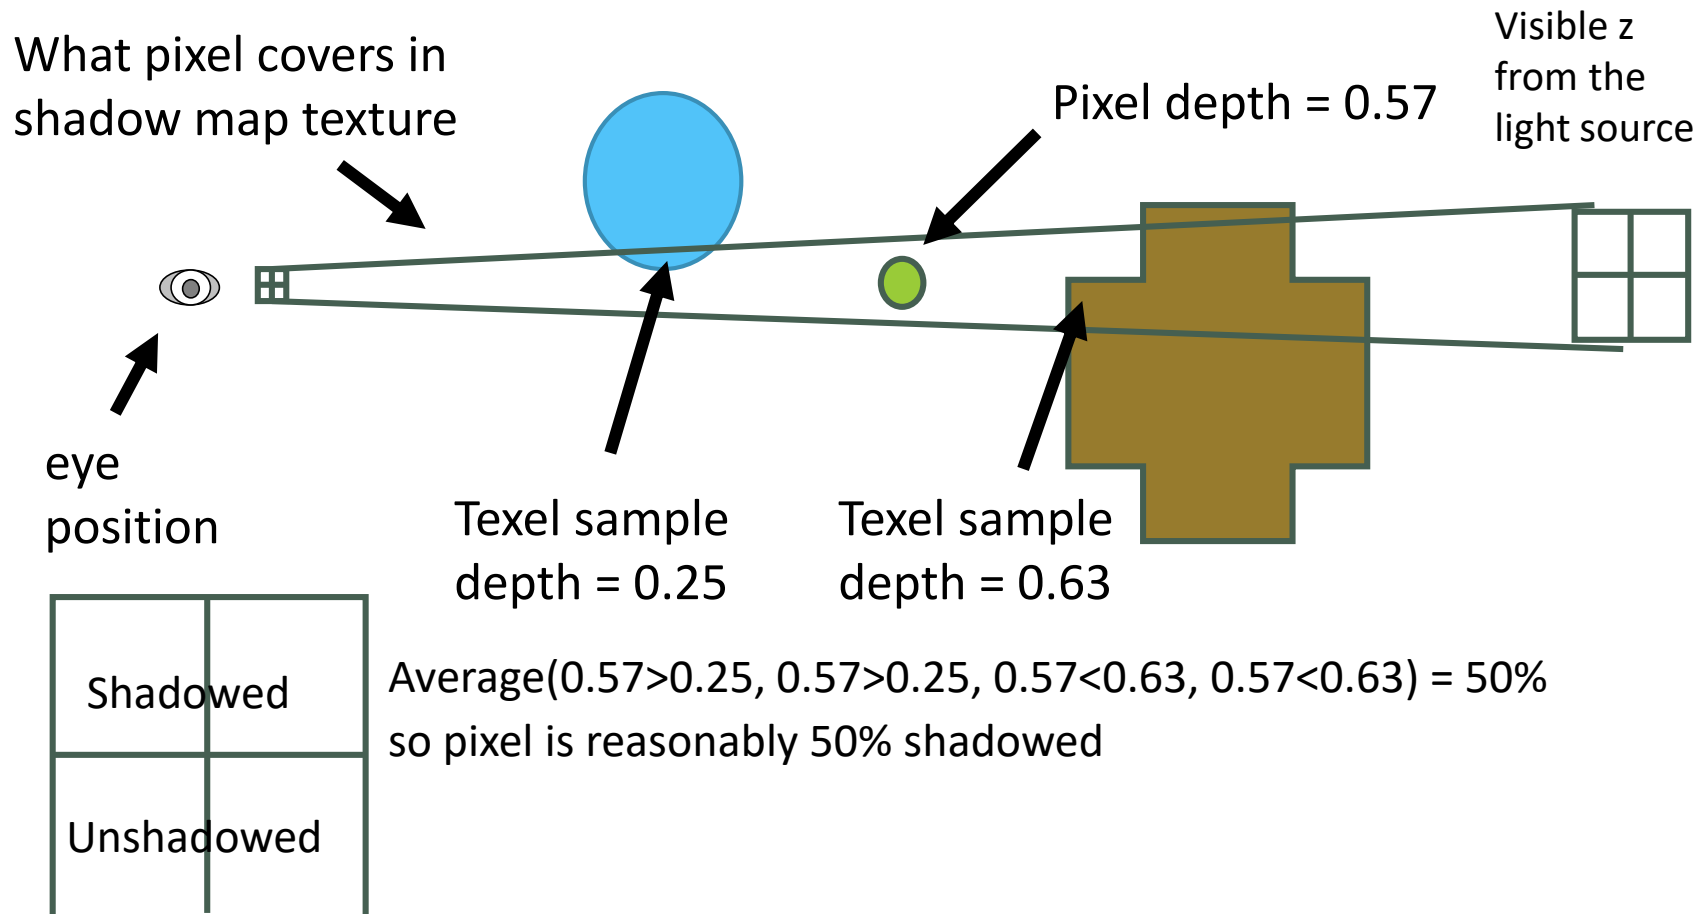
What pixel covers in shadow map texture



# Percentage Closer Filtering


Note: Alternative ways:  
Sampling surrounding  
texels of the depth map of  
light.

- Average comparison *results*, not depth values




# Percentage Closer Filtering

- ▶ Using a bigger filter produces fake soft shadows.




# Shadow volume


- ▶ Instead of image-space shadow mapping, shadow volume takes scene geometry into account.
- ▶ Check whether a polygon is within the shadow “volume”.
- ▶ Counts the number of intersections of “shadow volume polygons”.



# Stencil Enter/Leave Counting Approach




# Stencil Enter/Leave Counting Approach



$$\text{Shadow Volume Count} = +1 + 1 + 1 - 1 - 1 - 1 = 0$$


From C.Everitt and M.J.Kilgard, Practical & Robust Stenciled Shadow Volumes for Hardware-Accelerated Rendering, Nvidia

# Stencil Enter/Leave Counting Approach



**Shadow Volume Count =  $+1+1+1-1 = 2$**

# Stencil Enter/Leave Counting Approach



**Shadow Volume Count = 0**




# The stencil buffer


- ▶ Used by the shadow volume algorithm
  - ▶ often 8 bits per pixel.
- ▶ When rendering to it, we can add, subtract, etc.
- ▶ Then, the resulting image can be used to mask off subsequent rendering.
  - ▶ Can specify different rendering operations for each case:
    - ▶ stencil test fails
    - ▶ stencil test passes & depth test fails
    - ▶ stencil test passes & depth test passes

# Stencil buffer – real-time mirror

- ▶ Clear frame, depth & stencil buffers
- ▶ Draw all non-mirror geometry to frame & depth buffers
- ▶ Draw mirror to stencil buffer, where depth buffer passes
- ▶ Set depth to infinity, where stencil buffer passes
- ▶ Draw reflected geometry to frame & depth buffer, where stencil buffer passes



without stencil buffer



Reflected objects




with stencil buffer


# The shadow volume algorithm (Z pass)

- ▶ Rendered in four passes:
  - ▶ Pass 1: Draw the scene using ambient lighting only.
    - ▶ This is to setup the Z-buffer.
    - ▶ Turn off z-buffer and color buffer updates
    - ▶ (In later operations, only write to **stencil buffer**)
  - ▶ Pass 2: Draw front-facing shadow volume polygons.
    - ▶ Set stencil buffer to increment if depth test passes
  - ▶ Pass 3: Draw back-facing shadow volume polygons.
    - ▶ Set stencil buffer to decrement if depth test passes
  - ▶ Pass 4: Draw diffuse and specular lighting.
    - ▶ Set stencil buffer to mask all pixels where stencil buffer value is not 0.


# Stencil buffer for Z-pass algorithm




without shadow




with shadow




+




+



=



# Problems created by near plane clipping (Zpass approach)



# Shadow map vs. shadow volume

## ► Shadow Volumes

- Good: Anything can shadow anything, including self-shadowing, and the shadows are delicate.
- Bad: 3 or 4 passes, shadow volume polygons must be generated and rendered (lots of polygons & fill), intensive computation.
- Ugly: frustum problems for z-pass.

## ► Shadow Maps

- Good: Anything to anything, constant cost regardless of scene complexity, map can sometimes be reused.
- Bad: Frustum limited.
- Ugly: Jagged shadows if resolution is too low, biasing headaches.

**The End of Chapter**